

Speed Up Your Data Processing

Parallel and Asynchronous Programming in Data Science

By: Chin Hwee Ong (@ongchinhwee)

PyCon Taiwan 2020
6 September 2020

About me

Ong Chin Hwee 王敬惠

- Data Engineer @ ST Engineering
- Background in aerospace engineering + computational modelling
- Contributor to pandas
- Mentor team at BigDataX



@ongchinhwee

A typical data science workflow

1. Extract raw data
2. Process data
3. Train model
4. Evaluate and deploy model

Bottlenecks in a data science project

- Lack of data / Poor quality data
- Data processing
 - The 80/20 data science dilemma
 - In reality, it's closer to 90/10

Data Processing in Python

- **For loops in Python**

- Run on the **interpreter**, not compiled
- Slow compared with C

```
a_list = []  
for i in range(100):  
    a_list.append(i*i)
```

Data Processing in Python

- **List comprehensions**

- **Slightly faster** than for loops
- No need to call append function at each iteration

```
a_list = [i*i for i in range(100)]
```

Challenges with Data Processing

- **Pandas**

- Optimized for **in-memory analytics** using DataFrames
- **Performance + out-of-memory issues** when dealing with large datasets (> 1 GB)

```
import pandas as pd
import numpy as np
df = pd.DataFrame(list(range(100)))
squared_df = df.apply(np.square)
```

Challenges with Data Processing

- “Why not just use a Spark cluster?”

Communication overhead: Distributed computing involves communicating between **(independent) machines across a network!**

“Small Big Data”(*): Data too big to fit in memory, but not large enough to justify using a Spark cluster.

(*) Inspired by “The Small Big Data Manifesto”. Itamar Turner-Trauring (@itamarst) gave a great talk about Small Big Data at PyCon 2020.

@ongchinwee

What is parallel processing?

Let's imagine I work at a *kopi tiam* (咖啡店).



@ongchinhwee

Task 1: Toast 100 slices of bread

Assumptions:

1. I'm using single-slice toasters.
(Yes, they actually exist.)
2. Each slice of toast takes 2 minutes to make.
3. No overhead time.


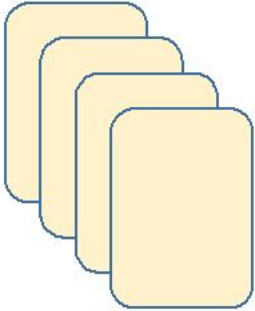
Image taken from:

<https://www.mitsubishielectric.co.jp/home/breadoven/product/to-st1-t/feature/index.html>





@ongchinhwee

Sequential Processing



= 25 bread slices

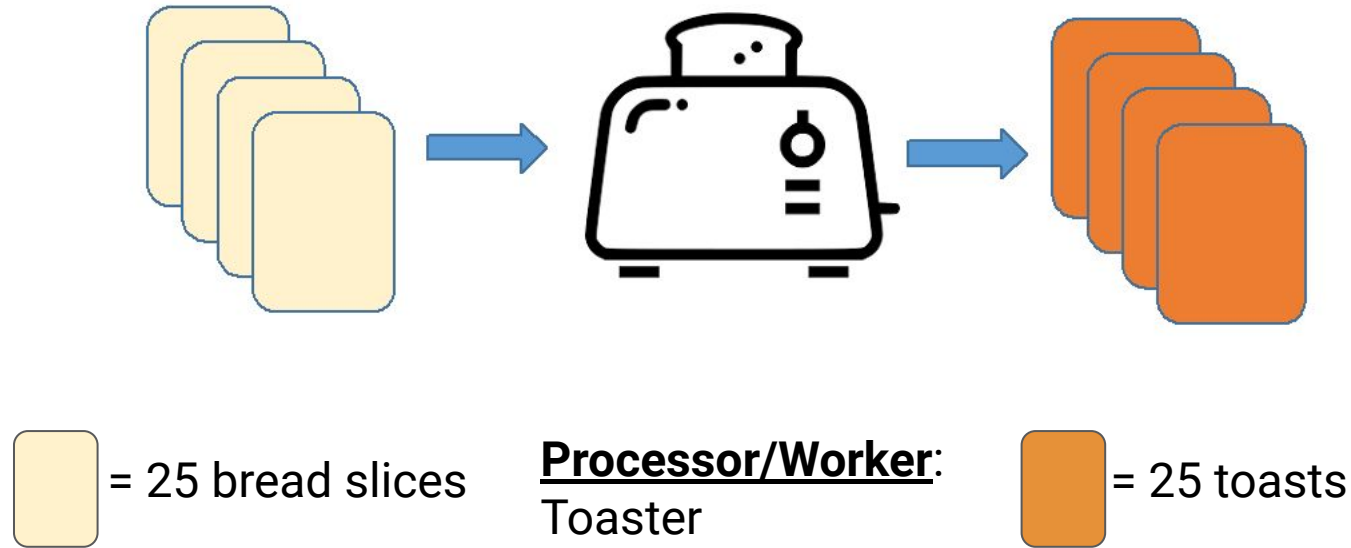
Sequential Processing



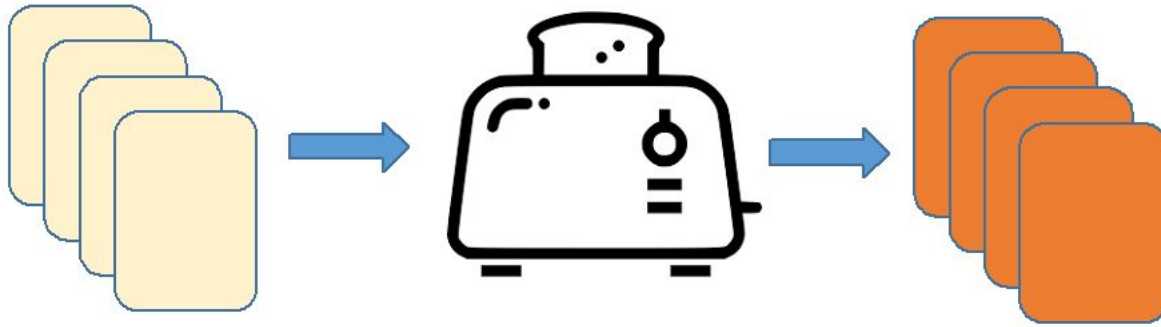
= 25 bread slices

Processor/Worker:
Toaster

Sequential Processing





Sequential Processing



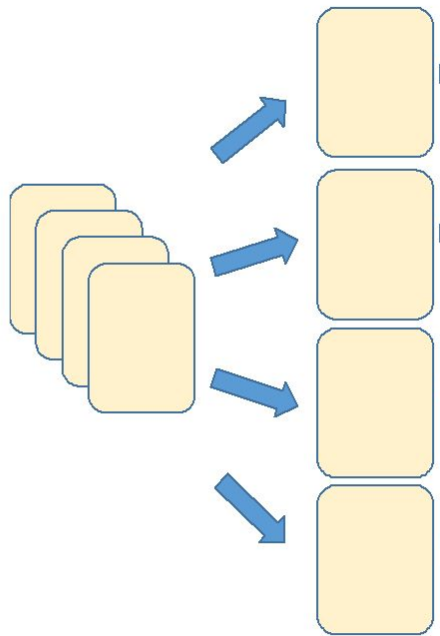
Execution Time = 100 toasts × 2 minutes/toast
= **200 minutes**

Parallel Processing

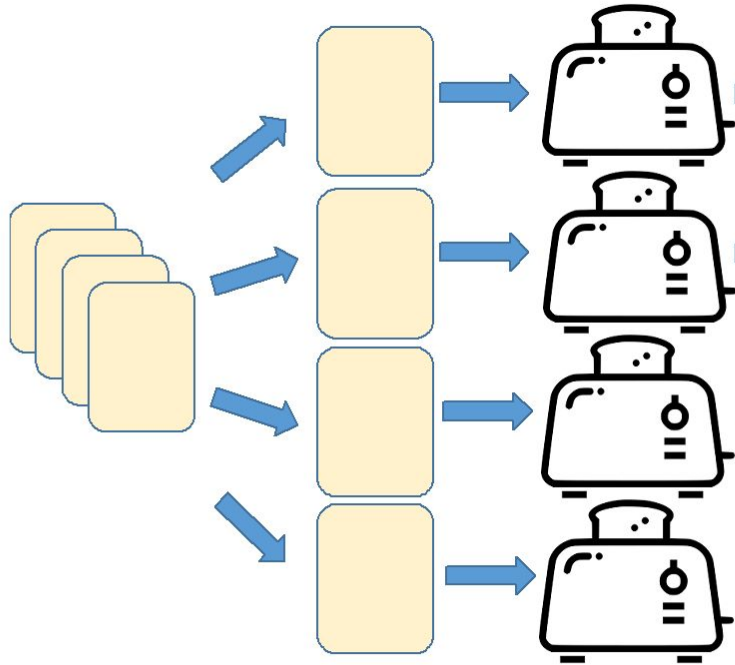


 = 25 bread slices

Parallel Processing

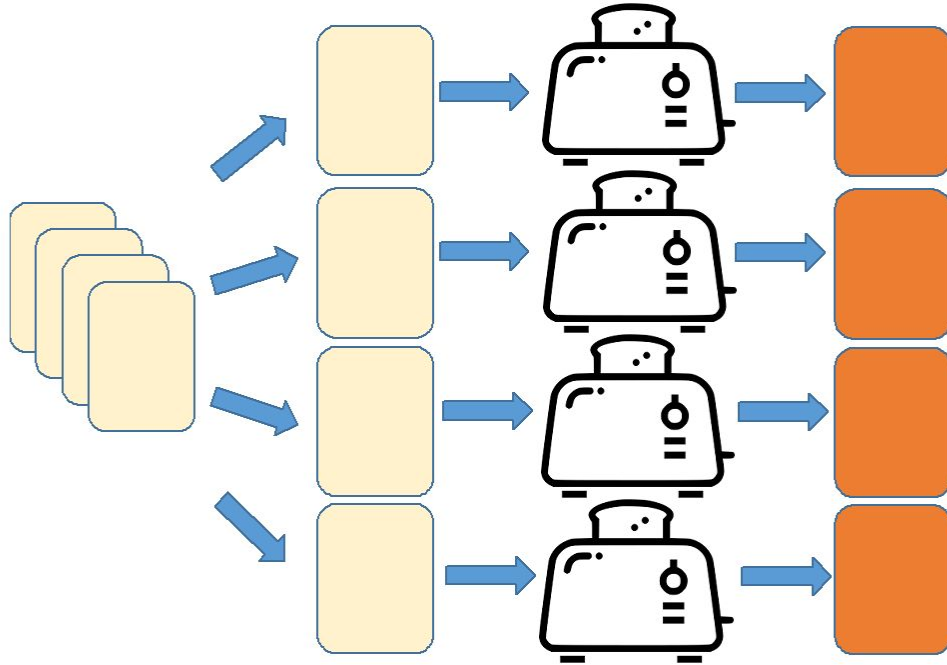


Parallel Processing



Processor (Core):
Toaster

Parallel Processing

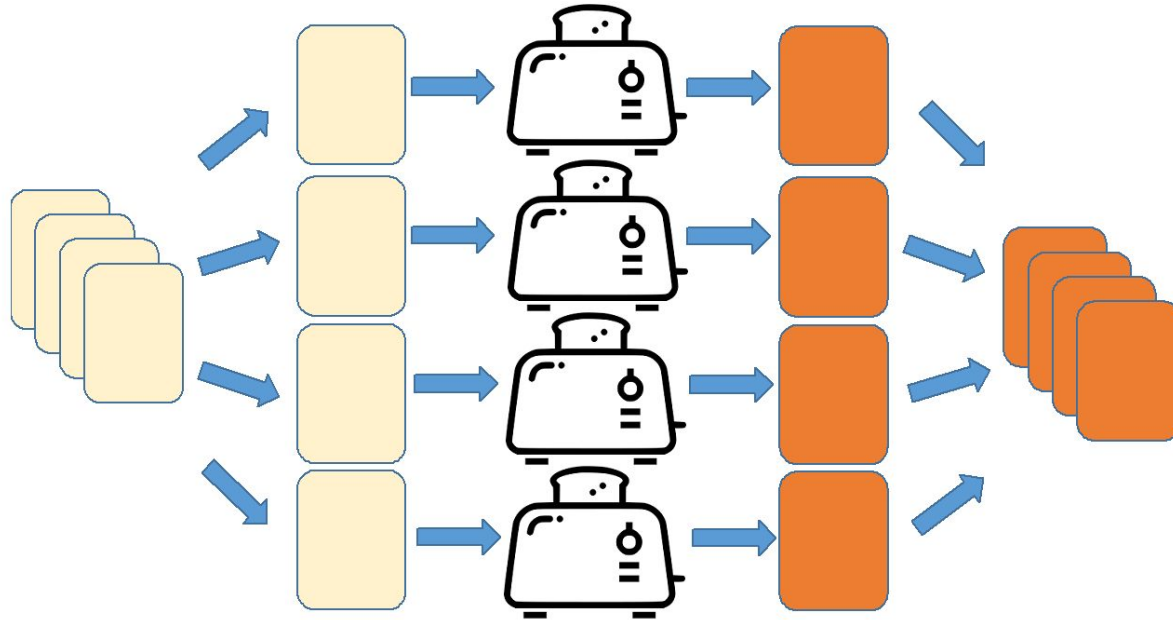


Processor (Core):
Toaster

Task is executed using a **pool** of **4 toaster subprocesses**.

Each toasting subprocess runs **in parallel** and **independently** from each other.


Parallel Processing



Processor (Core):
Toaster

Output of each
toasting process is
consolidated and
returned as an overall
output (which may or
may not be ordered).

Parallel Processing



Execution Time
= $100 \text{ toasts} \times 2 \text{ minutes/toast} \div 4 \text{ toasters}$
= **50 minutes**

Speedup
= **4 times**

Synchronous vs Asynchronous Execution

What do you mean by “Asynchronous”?

Task 2: Brew coffee

Assumptions:

1. I can do other stuff while making coffee.
2. One coffee maker to make one cup of coffee.
3. Each cup of coffee takes 5 minutes to make.



Image taken from: <https://mothership.sg/2020/08/kopimatic-machine/>

@ongchinhwee

Synchronous Execution



Task 2: Brew a cup of coffee on
coffee machine
Duration: 5 minutes

Synchronous Execution



Task 1: Toast two slices of bread on single-slice toaster after Task 2 is completed
Duration: 4 minutes

Task 2: Brew a cup of coffee on coffee machine
Duration: 5 minutes



Synchronous Execution



Task 1: Toast two slices of bread on single-slice toaster after Task 2 is completed
Duration: 4 minutes

Task 2: Brew a cup of coffee on coffee machine
Duration: 5 minutes



Output: **2 toasts + 1 coffee**

Total Execution Time = 5 minutes + 4 minutes = **9 minutes**


@ongchinhwee

Asynchronous Execution

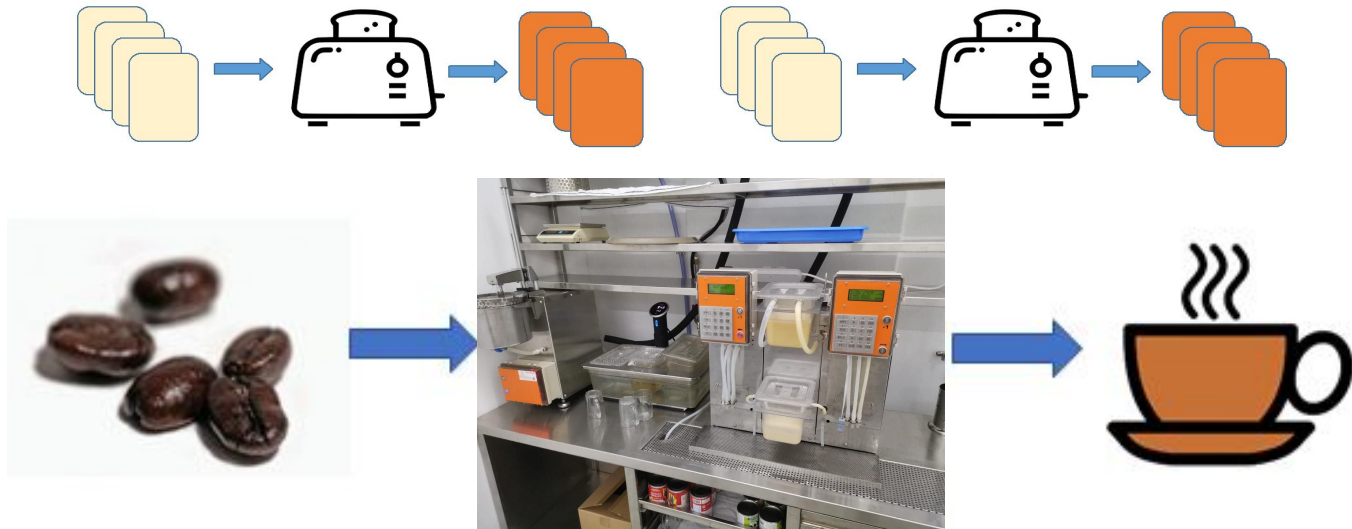
While brewing coffee:



Make some toasts:



Asynchronous Execution



Output: 2 toasts + 1 coffee
Total Execution Time = 5 minutes

When is it a good idea to go for parallelism?

(or, “Is it a good idea to simply buy a 256-core processor and parallelize all your codes?”)

Practical Considerations

- Is your code already optimized?
 - Sometimes, you might need to rethink your approach.
 - Example: Use list comprehensions or map functions instead of for-loops for array iterations.

Practical Considerations

- Is your code already optimized?
- Problem architecture
 - Nature of problem limits how successful parallelization can be.
 - If your problem consists of processes which depend on each others' outputs (Data dependency) and/or intermediate results (Task dependency), maybe not.

Practical Considerations

- Is your code already optimized?
- Problem architecture
- Overhead in parallelism
 - There will always be parts of the work that cannot be parallelized. → **Amdahl's Law**
 - Extra time required for coding and debugging (parallelism vs sequential code) → **Increased complexity**
 - **System overhead** including **communication overhead**

Amdahl's Law and Parallelism

Amdahl's Law states that the theoretical speedup is defined by the fraction of code p that can be parallelized:

$$S = \frac{1}{(1 - p) + \frac{p}{N}}$$


S : Theoretical speedup (theoretical latency)

p : Fraction of the code that can be parallelized

N : Number of processors (cores)

Amdahl's Law and Parallelism


If there are no parallel parts ($p = 0$): Speedup = 0



Amdahl's Law and Parallelism

If there are no parallel parts ($p = 0$): Speedup = 0

If all parts are parallel ($p = 1$):
Speedup = $N \rightarrow \infty$

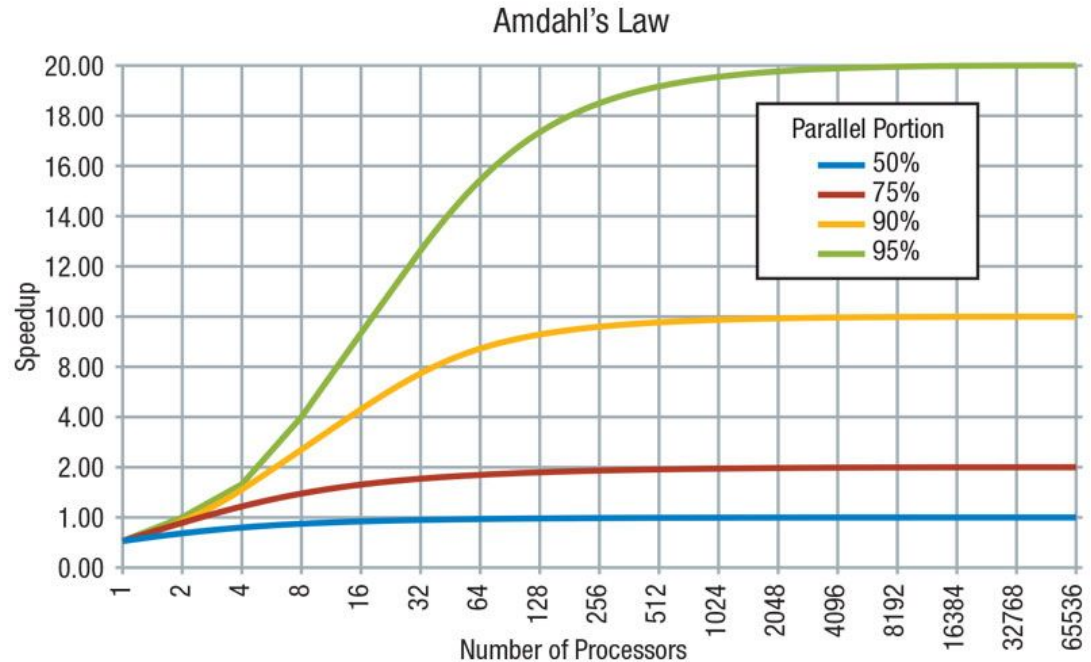


Amdahl's Law and Parallelism

If there are **no parallel parts** ($p = 0$): **Speedup = 0**

If **all parts are parallel** ($p = 1$):
Speedup = $N \rightarrow \infty$

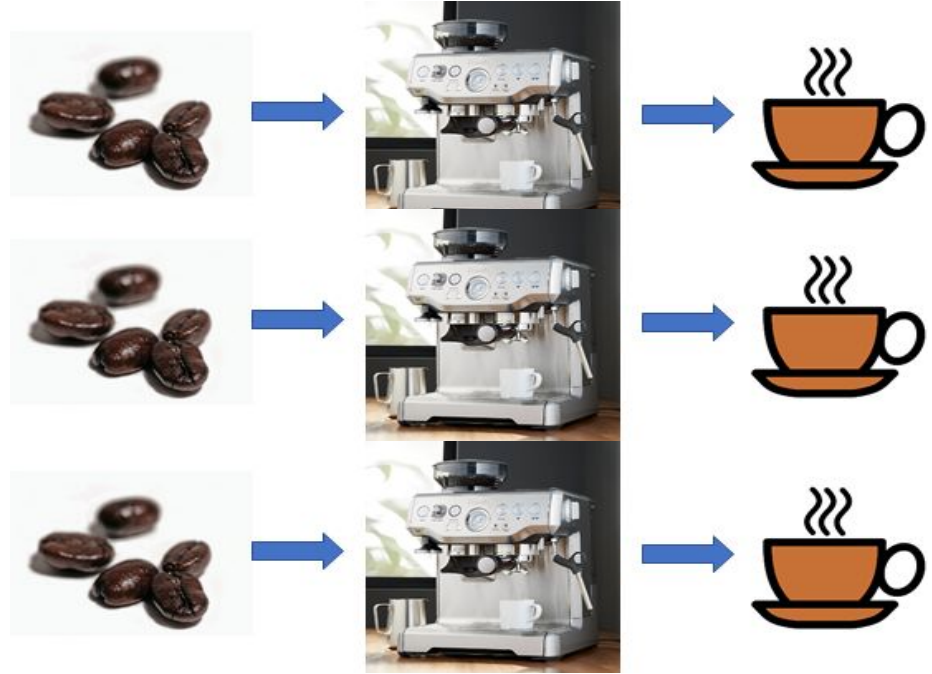
Speedup is limited by **fraction of the work that is not parallelizable** - will not improve even with infinite number of processors



Multiprocessing vs Multithreading

Multiprocessing:

System allows executing multiple processes at the same time using multiple processors



Multiprocessing vs Multithreading

Multiprocessing:

System allows executing multiple processes at the same time using multiple processors

Multithreading:

System executes multiple threads of sub-processes at the same time within a single processor

Multiprocessing vs Multithreading

Multiprocessing:

System allows executing multiple processes at the same time using multiple processors

Better for **processing large volumes of data**

Multithreading:

System executes multiple threads of sub-processes at the same time within a single processor

Best suited for **I/O or blocking operations**

Some Considerations

Data processing tends to be **more compute-intensive**

→ Switching between threads become increasingly inefficient

→ **Global Interpreter Lock (GIL)** in Python does not allow parallel thread execution

Did some pythonic developer just say



THREADS?

How to do Parallel + Asynchronous in Python?

(for data processing workflows (**))

(**) Common machine-learning libraries (e.g. scikit-learn, Tensorflow) already have their own implementation of multiprocessing

@ongchinwee

Parallel + Asynchronous Programming in Python

concurrent.futures module

- High-level API for launching **asynchronous (async) parallel tasks**
- Introduced in Python 3.2 as an abstraction layer over **multiprocessing** module
- Two modes of execution:
 - *ThreadPoolExecutor()* for async multithreading
 - *ProcessPoolExecutor()* for async multiprocessing

ProcessPoolExecutor vs ThreadPoolExecutor

From the Python Standard Library documentation:

For *ProcessPoolExecutor*, this method chops iterables into a number of chunks which it submits to the pool as *separate tasks*. The (approximate) size of these chunks can be specified by *setting chunksize to a positive integer*. For *very long iterables*, using a large value for chunksize can significantly improve performance compared to the default size of 1. With *ThreadPoolExecutor*, chunksize has *no effect*.

ProcessPoolExecutor vs ThreadPoolExecutor

ProcessPoolExecutor:

System allows executing multiple processes asynchronously using multiple processors

Uses multiprocessing module - side-steps GIL

ThreadPoolExecutor:

System executes multiple threads of sub-processes asynchronously within a single processor

Subject to GIL - not truly "concurrent"

@ongchinhwee

submit() in concurrent.futures

Executor.submit() takes as input:

1. The function (callable) that you would like to run, and
2. Input arguments (*args, **kwargs) for that function;

and returns a futures object that **represents the execution of the function.**

map() in concurrent.futures

Similar to map(), **Executor.map()** takes as input:

1. The function (callable) that you would like to run, and
2. A list (iterable) where each element of the list is a single input to that function;

and returns an iterator that **yields** the results of the function being applied to every element of the list.

Case: Network I/O Operations

Dataset: Data.gov.sg Realtime Weather Readings
(<https://data.gov.sg/dataset/realtime-weather-readings>)

API Endpoint URL: <https://api.data.gov.sg/v1/environment/>

Response: JSON format

Initialize Python modules

```
import numpy as np
```

```
import requests  
import json
```

```
import sys  
import time  
import datetime  
from tqdm import trange, tqdm  
from time import sleep  
from retrying import retry  
import threading
```

Initialize API request task

```
@retry(wait_exponential_multiplier=1000, wait_exponential_max=10000)
def get_airtemp_data_from_date(date):
    print('{}: running {}'.format(threading.current_thread().name,
                                   date))
    # for daily API request
    url =
    "https://api.data.gov.sg/v1/environment/air-temperature?date="\
        + str(date)
    JSONContent = requests.get(url).json()
    content = json.dumps(JSONContent, sort_keys=True)
    sleep(1)
    print('{}: done with {}'.format(
        threading.current_thread().name, date))
    return content
```

threading module to
monitor thread
execution

@ongchinhwee

Initialize Submission List

```
date_range = np.array(sorted(  
    [datetime.datetime.strftime(  
        datetime.datetime.now() - datetime.timedelta(i)  
, '%Y-%m-%d') for i in range(100)]))
```

Using List Comprehensions

```
start_cpu_time = time.clock()
```

```
data_np = [get_airtemp_data_from_date(str(date)) for date in  
tqdm(date_range)]
```

```
end_cpu_time = time.clock()
```

```
print(end_cpu_time - start_cpu_time)
```

Using List Comprehensions

```
start_cpu_time = time.clock()
```

List Comprehensions:

977.88 seconds (~ 16.3mins)

```
data_np = [get_airtemp_data_from_date(str(date)) for date in  
tqdm(date_range)]
```

```
end_cpu_time = time.clock()
```

```
print(end_cpu_time - start_cpu_time)
```

Using ThreadPoolExecutor

```
from concurrent.futures import ThreadPoolExecutor, as_completed
```

```
start_cpu_time = time.clock()
```

```
with ThreadPoolExecutor() as executor:  
    future = {executor.submit(get_airtemp_data_from_date, date):date  
              for date in tqdm(date_range)}  
resultarray_np = [x.result() for x in as_completed(future)]
```

```
end_cpu_time = time.clock()
```

```
total_tpe_time = end_cpu_time - start_cpu_time
```

```
sys.stdout.write('Using ThreadPoolExecutor: {} seconds.\n'.format(  
    total_tpe_time))
```

@ongchinhwee

Using ThreadPoolExecutor

```
from concurrent.futures import ThreadPoolExecutor, as_completed
```

```
start_cpu_time = time.clock()
```

ThreadPoolExecutor (40 threads):
46.83 seconds (~20.9 times faster)

```
with ThreadPoolExecutor() as executor:  
    future = {executor.submit(get_airtemp_data_from_date, date):date  
              for date in tqdm(date_range)}  
resultarray_np = [x.result() for x in as_completed(future)]
```

```
end_cpu_time = time.clock()
```

```
total_tpe_time = end_cpu_time - start_cpu_time
```

```
sys.stdout.write('Using ThreadPoolExecutor: {} seconds.\n'.format(  
    total_tpe_time))
```

@ongchinhwee

Case: Image Processing

Dataset: Chest X-Ray Images (Pneumonia)

(<https://www.kaggle.com/paultimothymooney/chest-xray-pneumonia>)

Size: 1.15GB of x-ray image files with normal and pneumonia (viral or bacterial) cases

Data Quality: Images in the dataset are of **different dimensions**

Initialize Python modules

```
import numpy as np  
from PIL import Image
```

```
import os  
import sys  
import time
```

Initialize image resize process

```
def image_resize(filepath):  
    '''Resize and reshape image'''  
    sys.stdout.write('{}: running {}\n'.format(os.getpid(), filepath))  
    im = Image.open(filepath)  
    resized_im = np.array(im.resize((64, 64)))  
    sys.stdout.write('{}: done with  
{ }\n'.format(os.getpid(), filepath))  
    return resized_im
```

os.getpid() to
monitor process
execution

Initialize File List in Directory

```
DIR = './chest_xray/train/NORMAL/'
```

No. of images in
'train/NORMAL': **1431**

```
train_normal = [DIR + name for name in os.listdir(DIR)  
                 if os.path.isfile(os.path.join(DIR, name))]
```

Using map()

```
start_cpu_time = time.clock()
```

```
result = map(image_resize, train_normal)
```

```
output = np.array([x for x in result])
```

```
end_cpu_time = time.clock()
```

```
total_tpe_time = end_cpu_time - start_cpu_time
```

```
sys.stdout.write('Map completed in {}
```

```
seconds.\n'.format(total_tpe_time))
```

Using map()

```
start_cpu_time = time.clock()
```

```
result = map(image_resize, train_normal)
```

```
output = np.array([x for x in result])
```

```
end_cpu_time = time.clock()
```

```
total_tpe_time = end_cpu_time - start_cpu_time
```

```
sys.stdout.write('Map completed in {}
```

```
seconds.\n'.format(total_tpe_time))
```

map():

29.48 seconds

Using List Comprehensions

```
start_cpu_time = time.clock()
```

```
listcomp_output = np.array([image_resize(x) for x in  
train_normal])
```

```
end_cpu_time = time.clock()
```

```
total_tpe_time = end_cpu_time - start_cpu_time
```

```
sys.stdout.write('List comprehension completed in {}  
seconds.\n'.format(  
    total_tpe_time))
```

Using List Comprehensions

```
start_cpu_time = time.clock()
```

List Comprehensions:
29.71 seconds

```
listcomp_output = np.array([image_resize(x) for x in  
train_normal])
```

```
end_cpu_time = time.clock()  
total_tpe_time = end_cpu_time - start_cpu_time  
sys.stdout.write('List comprehension completed in {}  
seconds.\n'.format(  
    total_tpe_time))
```


Using ProcessPoolExecutor

```
from concurrent.futures import ProcessPoolExecutor  
start_cpu_time = time.clock()
```

```
with ProcessPoolExecutor() as executor:  
    future = executor.map(image_resize, train_normal)  
  
array_np = np.array([x for x in future])
```

```
end_cpu_time = time.clock()  
total_tpe_time = end_cpu_time - start_cpu_time  
sys.stdout.write('ProcessPoolExecutor completed in {}  
seconds.\n'.format(  
    total_tpe_time))
```

Using ProcessPoolExecutor

```
from concurrent.futures import ProcessPoolExecutor
```

```
start_cpu_time = time.clock()
```

ProcessPoolExecutor (8 cores):
6.98 seconds (~4.3 times faster)

```
with ProcessPoolExecutor() as executor:
```

```
    future = executor.map(image_resize, train_normal)
```

```
array_np = np.array([x for x in future])
```

```
end_cpu_time = time.clock()
```

```
total_tpe_time = end_cpu_time - start_cpu_time
```

```
sys.stdout.write('ProcessPoolExecutor completed in {}
```

```
seconds.\n'.format(
```

```
    total_tpe_time))
```

Key Takeaways

Not all processes should be parallelized

- Parallel processes come with **overheads**
 - Amdahl's Law on parallelism
 - System overhead including **communication overhead**
 - If the cost of rewriting your code for parallelization outweighs the time savings from parallelizing your code, consider **other ways of optimizing your code** instead.

Reach out to me!



: ongchinhwee



: @ongchinhwee



: hweecat



: <https://ongchinhwee.me>

And check out my slides on:



hweecat/talk_parallel-async-python

@ongchinhwee